Mohd Faiz

portfolio-mohd-faiz-19.vercel.app
github.com/Rexter9
linkedin.com/in/mohd-faiz-0b8773373
+91-6394544068
mhdfaiz1912@gmail.com

Professional Summary

Software Engineer with strong **backend**, **full-stack** and **AI** skills. Developed **RESTful APIs**, **JWT authentication**, **role-based access control (RBAC)**, **workflow automation** and **real-time** systems using **Node.js**, **Express.js**, **MongoDB**, **React.js** and **Python**. Built **Generative AI** and **agentic AI** systems using **RAG**, **LangChain**, **LangGraph** and **MCP**, backed by solid **DSA**, **OOP**, **DBMS** and **system design** fundamentals.

Education

Chandigarh University — **Mohali, Punjab**
B.E. in Computer Science — ***2021 – 2025***

Technical Skills

Languages:	C++, JavaScript, Python
Backend:	Node.js, Express.js, RESTful APIs, JWT Authentication, BullMQ, Apache Kafka
Databases:	MongoDB, Redis, PostgreSQL, Vector Databases
Generative AI:	LangChain, LangGraph, RAG, MCP
Frontend:	React.js, Next.js, Tailwind CSS
Dev Tools:	Docker, Git/GitHub, Vercel, Render, Streamlit
Core CS:	Data Structures & Algorithms, OOP, System Design, DBMS, Operating Systems, Computer Networks

Projects

Relay: AI Workflow Orchestrator | GitHub | Live Demo — ***Sep 2026 – Oct 2026***

Full-Stack AI Workflow Automation Platform

- Built an n8n-style **workflow orchestration engine** with a **React** console, where workflows are graphs of **AI calls**, **HTTP requests**, conditions, delays, notifications and **human approval gates** (**6 node types**). Shipped **5 ready-made templates** (refund processing, ticket routing, expense reimbursement and more) with **webhook** triggers secured by a shared secret.
- Designed asynchronous execution with **BullMQ + Redis** and a separate worker process. Persisted step-level state in **MongoDB** so interrupted runs **resume from the last completed step** and skip steps already recorded as completed.
- Added **prompt-injection checks** on **Gemini** inputs. **Dockerized** the stack into **4 containers** with Docker Compose and deployed the console on **Vercel**.

TaskFlow: Team Task Tracker API | GitHub — ***Jul 2026 – Aug 2026***

Backend API with Real-Time Collaboration

- Built a **Node.js, Express 5 and MongoDB** backend with **5 data models**, **JWT authentication** with bcrypt hashing and server-side token invalidation, and **role-based access** with **3 roles** (owner, admin, member) for teams, tasks, comments and attachments.
- Added **real-time notifications** over **Socket.io** with JWT-authenticated sockets and per-user and per-team rooms, plus task filtering, search, sorting and **pagination**.
- Integrated the **Claude API** to draft task titles and descriptions from short notes. Added **request validation**, **rate limiting**, **NoSQL-injection protection** and **Helmet** security headers, plus lint checks in **GitHub Actions (CI)**.

Resume Scoring JSON-RPC Server | GitHub — ***Jun 2026***

Async Python Server and Agent for Resume Scoring

- Built an asynchronous **Python** server with **aiohttp** exposing **13 methods** in **4 groups** over **JSON-RPC 2.0** (HTTP and WebSocket), including **directory watching** and **parallel batch processing**.
- Built a **5-step** agent that reads a job description and resumes only through the server's methods, scores each from 0 to 100 with **Groq LLaMA 3.3 70B** and saves ranked JSON. Covered by **20 pytest** tests and **9 error codes**.
- Used **async I/O** (aiohttp, aiofiles) with runtime method discovery through **rpc.discover** and configurable batch size and watch interval. Sample runs produce ranked results labelled **Strong, Moderate or Weak Match**.

HR Matching Agent | GitHub — ***May 2026***

LangGraph Based Resume Matching and Ranking Agent

- Built a **LangGraph** agent with a **6-node** flow (JD parsing, requirement extraction, semantic search, ranking, report generation, human feedback) using **ChromaDB**, HuggingFace embeddings and **Groq LLaMA 3.3 70B**.
- Added a **Streamlit** chat with **10 tools** for candidate comparison, match reports and tailored interview questions. Recruiters can change priorities mid-chat and the agent **re-ranks and explains the change** across **3 screening rounds**.
- Tracked **agent state** across turns (conversation history, parsed requirements, scored shortlist, screening round and refinement history) so each round builds on the last, and tested the conversation flow with **6 scenarios**.

Certifications

- **Web Design for Everybody: Basics of Web Development & Coding**, University of Michigan (Coursera)
- **Fundamentals of Visualization with Tableau**, University of California, Davis (Coursera)
- **Foundations of Project Management**, Google (Coursera)